

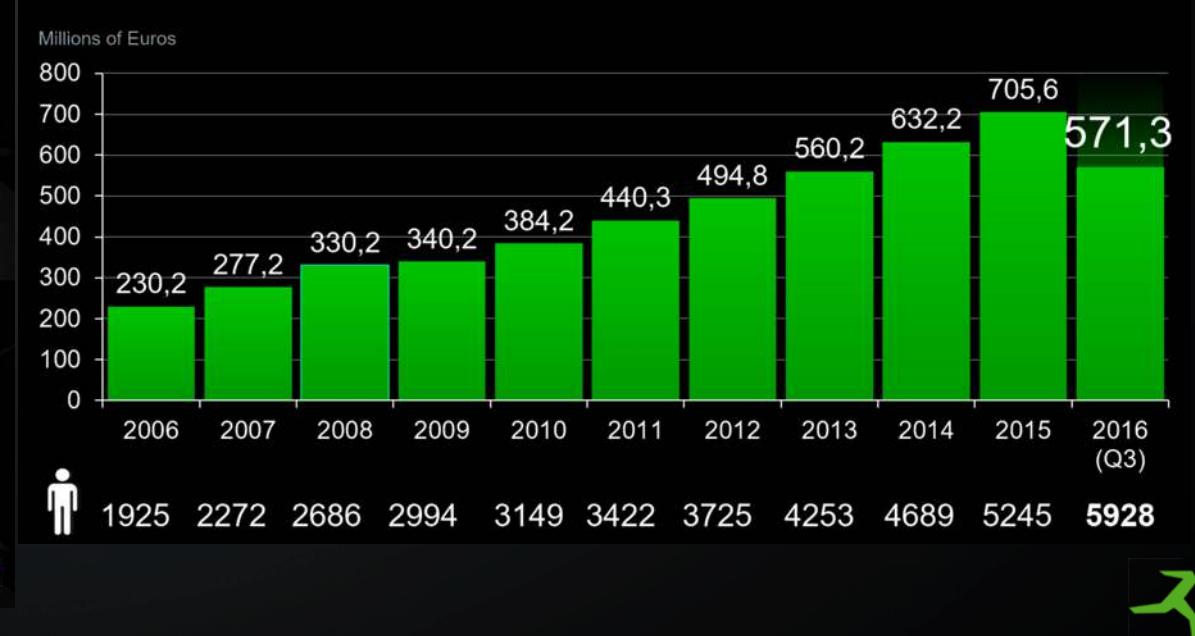
# Blockchain meets Machine Learning: How to develop your own intelligent smart contracts system.



Data Reply: Thomas Schmiedel [[t.schmiedel@reply.de](mailto:t.schmiedel@reply.de)]

# Short Intro: [Data] Reply

- Hi, my name is Thomas Schmiedel
- I'm data scientist at Data Reply in Munich, our specialty is Big Data, Data Science / Machine Learning and Data Engineering
- Data Reply is part of the Reply Network with over 6000 employees world-wide
- Strong focus on innovative IT-technologies



# Motivation



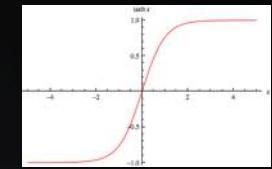
Why Blockchain **and** Machine Learning?

- Well, there's obviously some hype around each of those technologies, why not combine them and create even more hype? :-)
- The truth is: Every Blockchain network consists of decentralized participants, constantly exchanging messages between each other and keeping a record of history they all fully agree upon
- Up to now, the data inside that record is more or less unstructured, just a pile of bits and bytes
- Let's take it to the next level and build a colorful ecosystem on top of it, so people can collaborate more efficiently and find interesting things in seconds

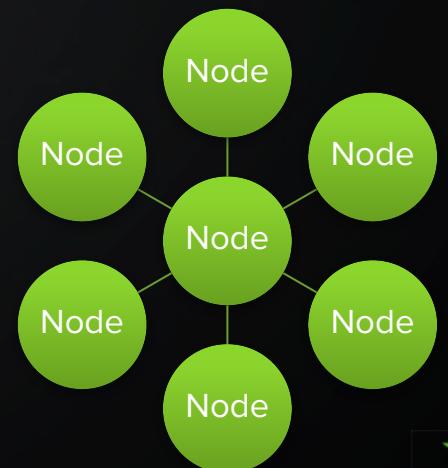


# Agenda

What are we going to do **now?**



- Well, at least some **theory** about Blockchains and Machine Learning is necessary, **we'll** do that first
- Then, **we'll** build an **architecture** for our intelligent Blockchain client
- Afterwards, **we'll** write a **smart contract** and deploy it to the network
- The **client software** must also be written
- Testing it, of course!

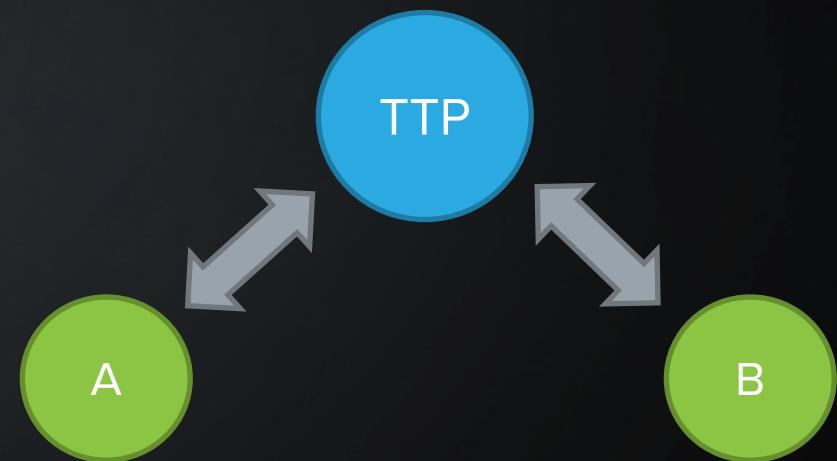


# Whom to Trust?

A short story about trust:



- Let's suppose you want to **buy a car** from a stranger, is there **trust involved**?
- Of course, he could show up with some “friends” and **drive away with car and money**
- Any idea what to do about that?
- Trusted Third Party (**TTP**):
  - Takes **car** from A and won't give it back
  - Receives **money** from B
  - Gives A the money and B the car
- Is there trust involved? Unfortunately yes, TTP could also run away
- You **can't eliminate trust by forwarding responsibility** to someone else (TTP, Server)



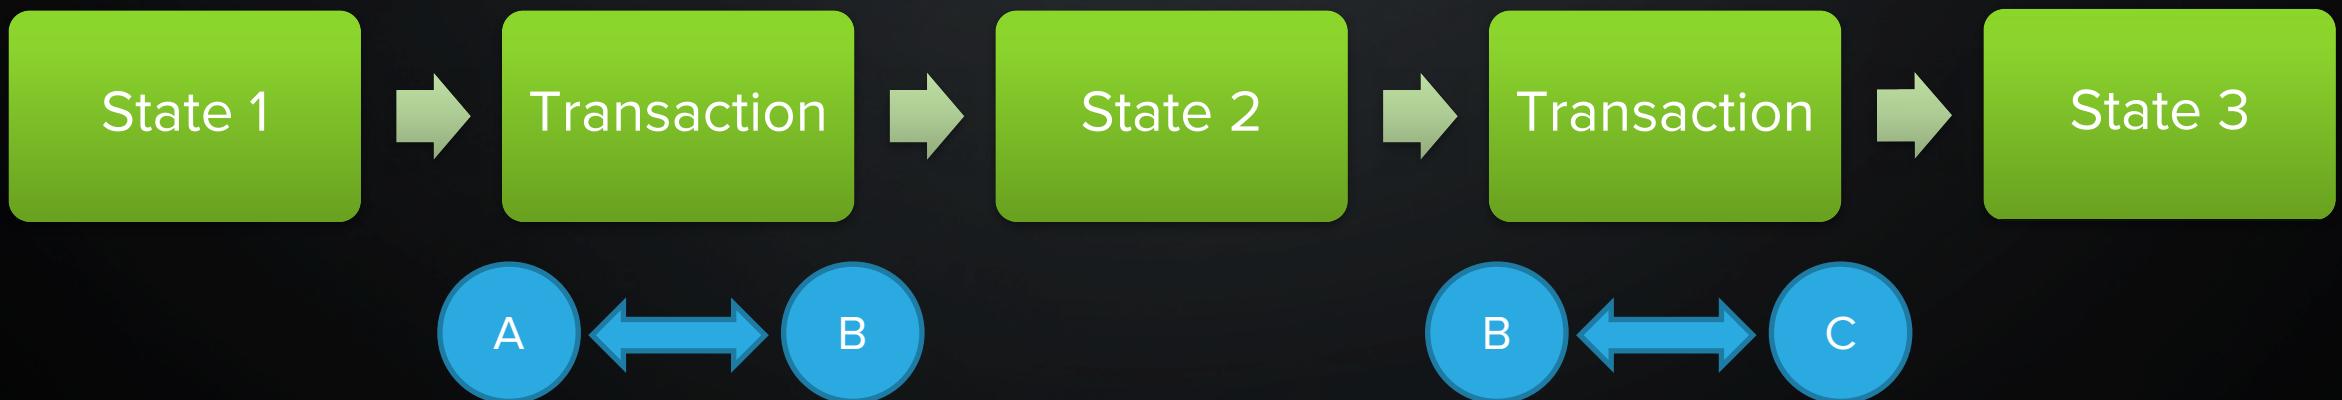
# Atomic Transactions

Solution:

- Atomic transactions: either car and money change ownership simultaneously or nothing happens

But how?

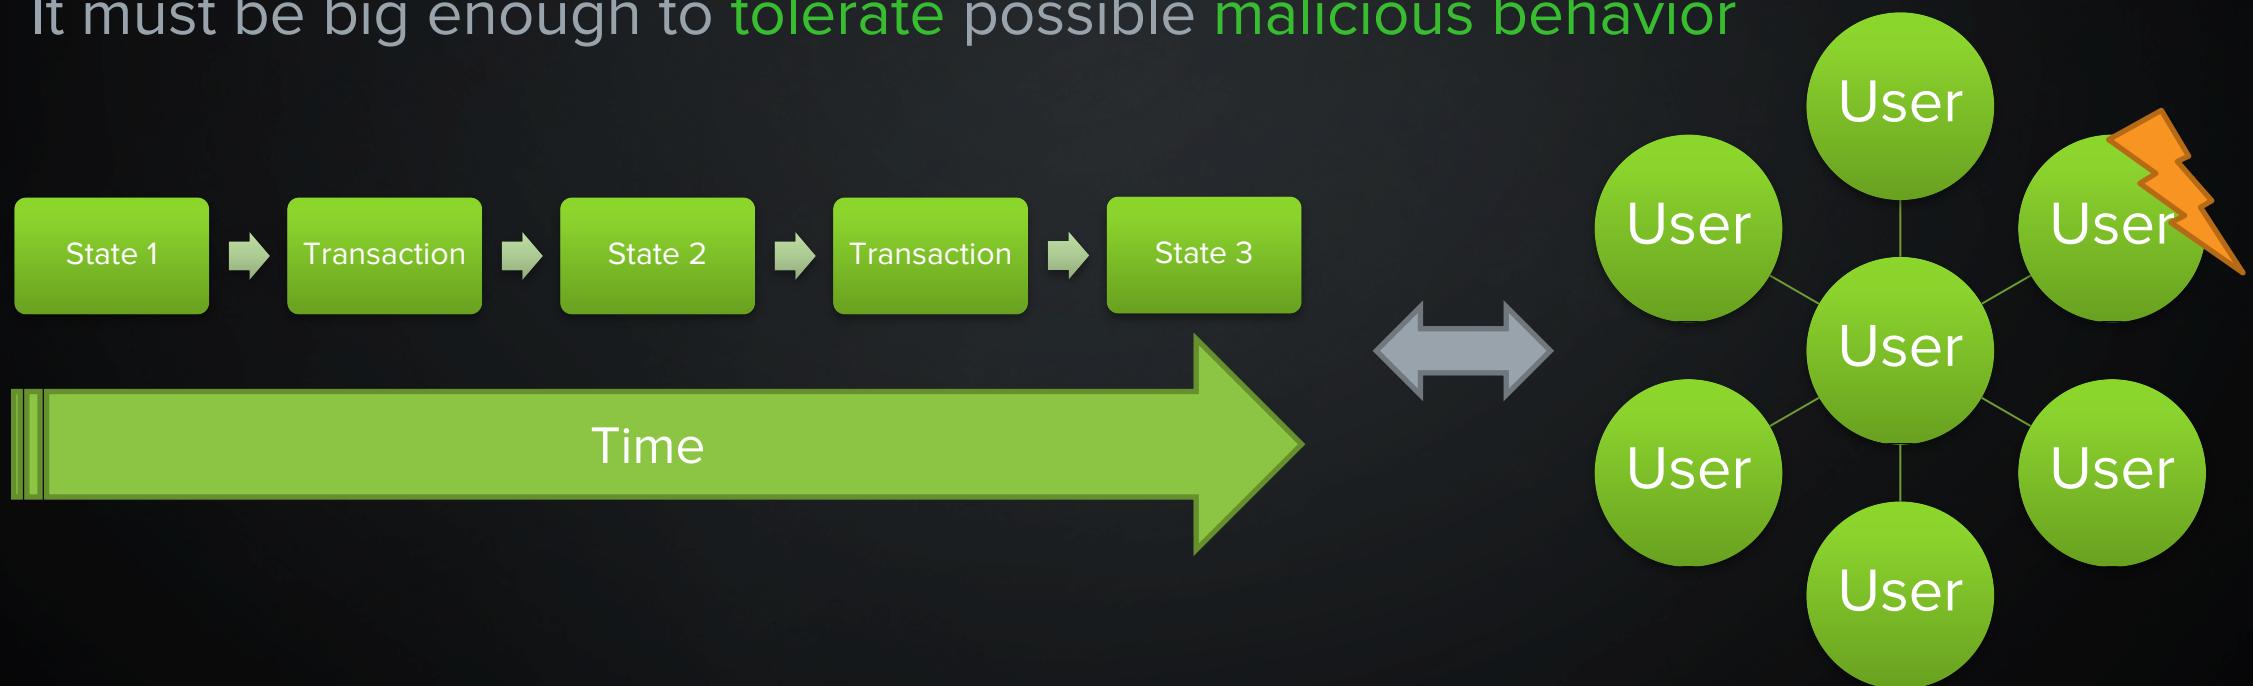
- Let's say we have some universal storage of ownership that everybody has access to
- And this storage can only be transformed according to certain rules
- Our rule: car and money can only change ownership in a single step
- → Our universal storage of ownership can only be altered by following the rule



# Core Feature: Decentralization

How to enforce following the rules?

- We need a whole network keeping the same record and verifying if changes to the record follow the rules
- It must be big enough to tolerate possible malicious behavior



# Closing the Circle

Now we introduce some different names:

- Universal storage of ownerships... let's call it Blockchain ✓
- The transaction rule... let's call it smart contract ✓
- The participants in our network... let's call them nodes or peers ✓

Therefore:

- Each node has an identical copy of the Blockchain and only allows transformations of it according to a smart contract
- But: nodes must always be in consensus (sync) about current state of reality
- → The complete record is transformed step by step from one consistent state to the next consistent state

How to keep nodes in sync? → Consensus Algorithm



# Keeping Consensus among Nodes

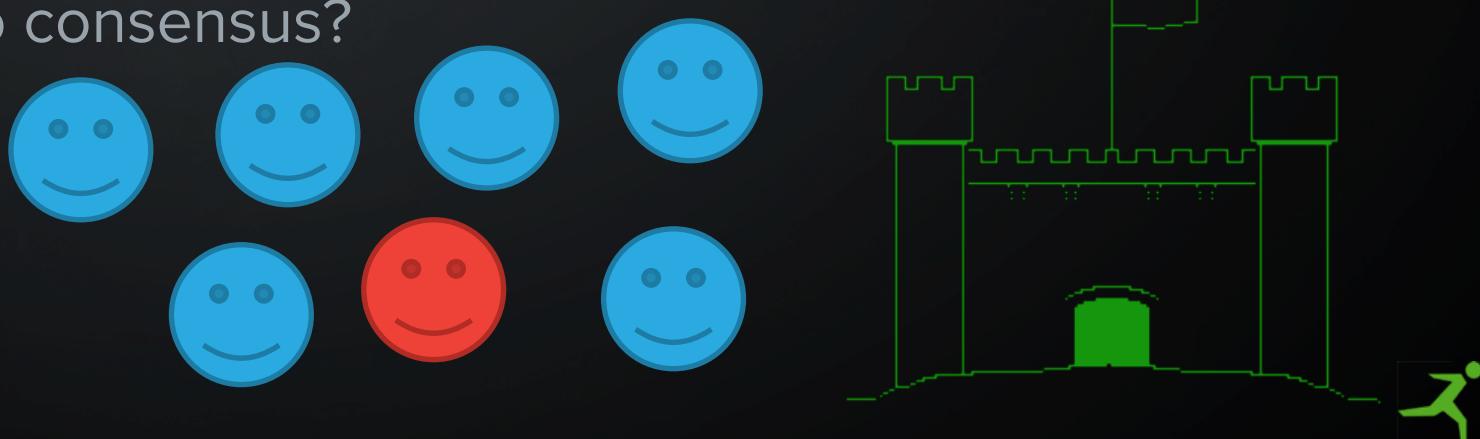
## Byzantine Fault Tolerance:

- Can the network **cope with malicious nodes** sending manipulated messages?
- How many malicious nodes are needed to cause **inconsistencies within the network state**? (some nodes assuming a different current reality than the rest)

## The Generals Problem:

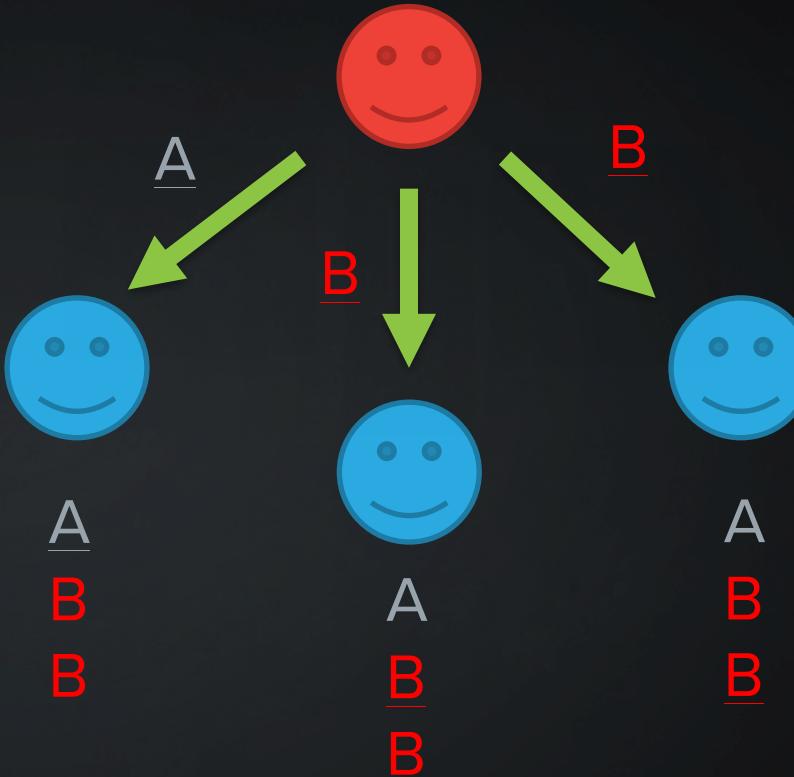
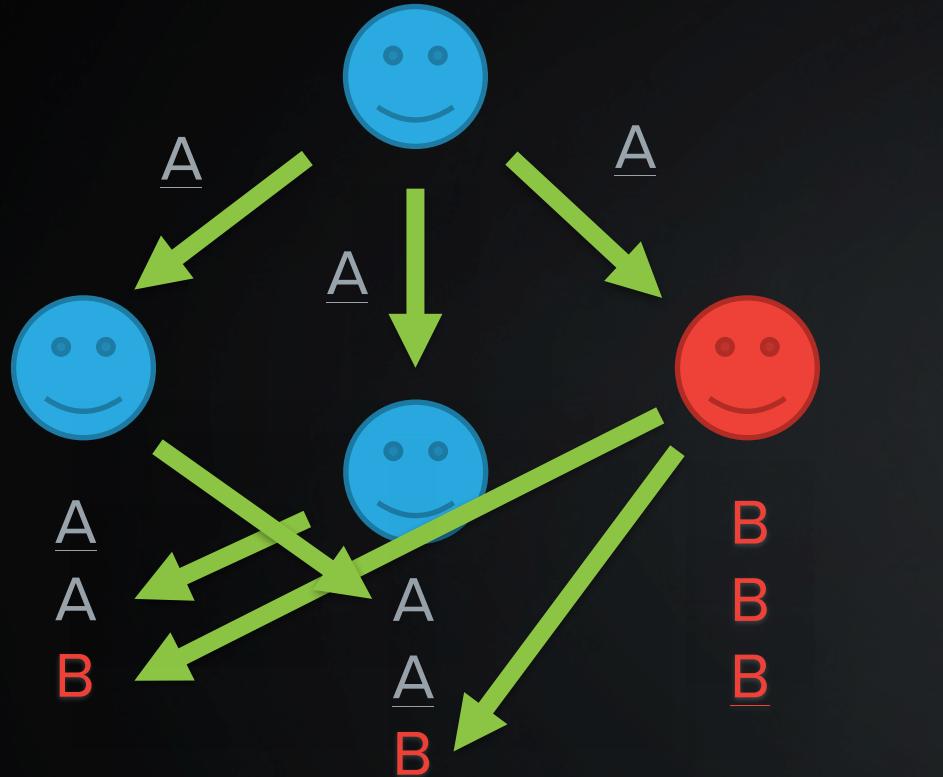
- N **Generals** want to attack a **fortress**, but they can only win if they **attack together**
- Among the generals there is a **traitor** (or more), trying to make the attack fail
- Suppose one of them gives a command, can you think of a mechanism to bring all remaining generals to consensus?

- Please try with pen and paper
- 5min

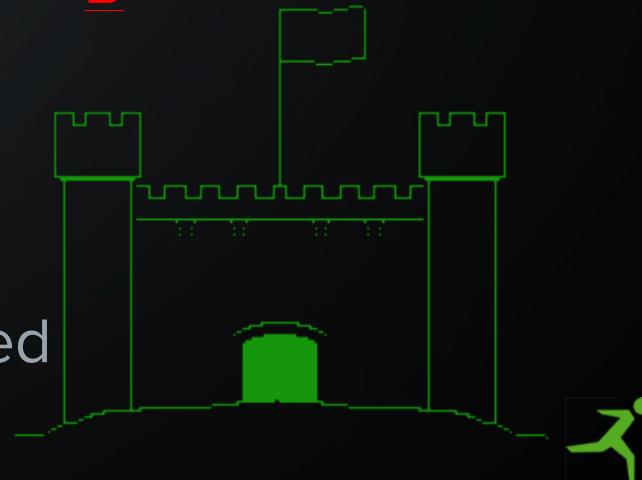


# Byzantine Fault Tolerance

See: Miguel Castro, Barbara Liskov: Practical Byzantine Fault Tolerance (1999)

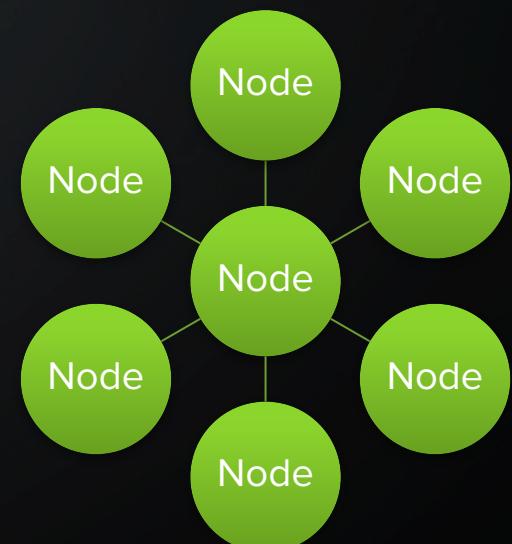
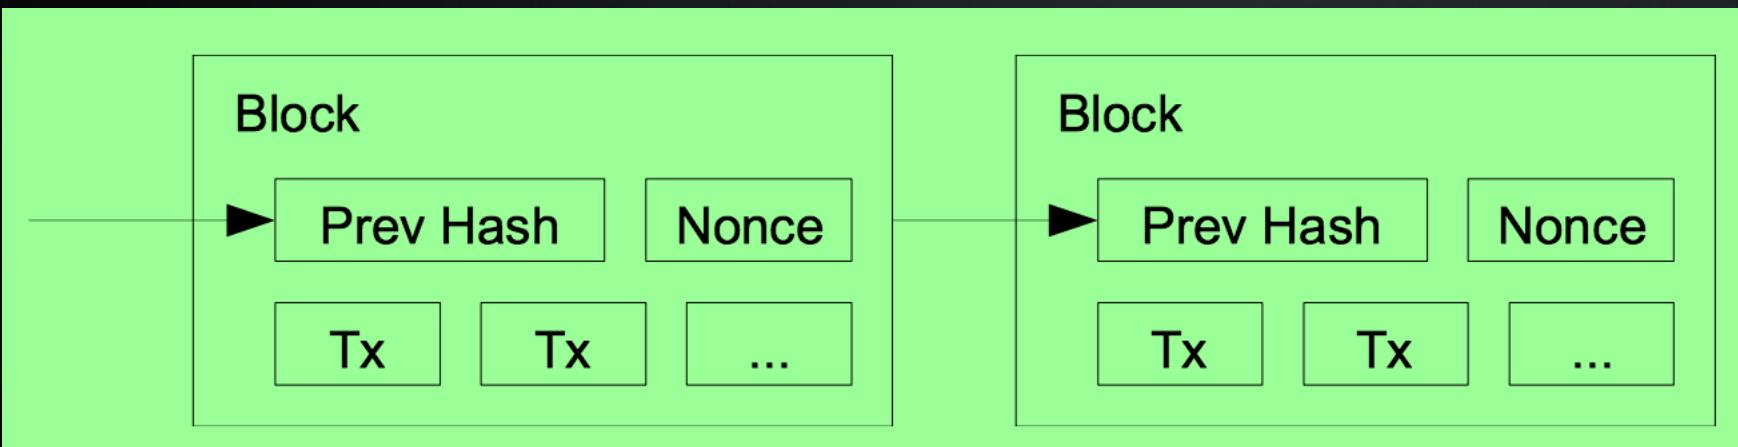


- One general sends a **request to all others**
- **They** again ask all others about what they received
- Each general does a **majority vote** on the responses he received
- To tolerate  $f$  traitors, we need  $n > 3f$  generals



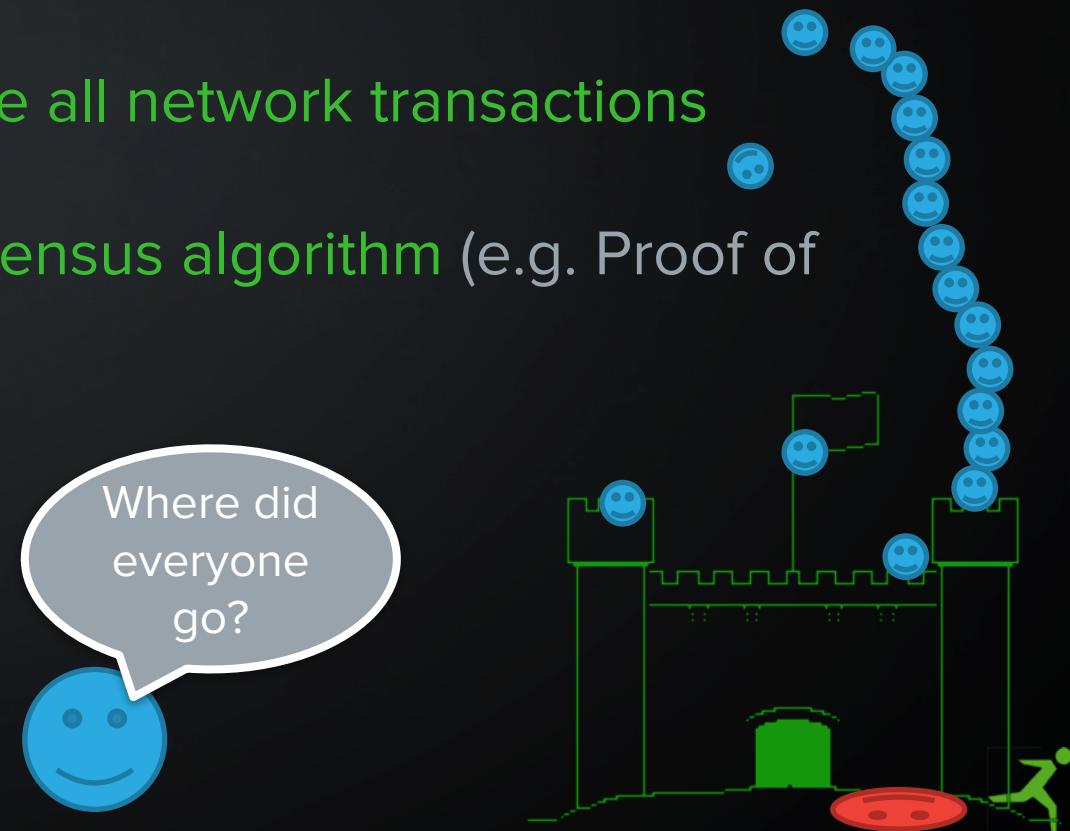
# Proof Of Work

- Different approach
- Some nodes (=miners) are computing sha256 hashes until the leading N bits of the resulting hash-sum are 0
- This ensures they do difficult computational work
- If a node finds a such a hash → new block
- The hash includes a hash of the block before and transactions that have occurred since the last block
- This forms a chain of blocks with payload (e.g. transactions)



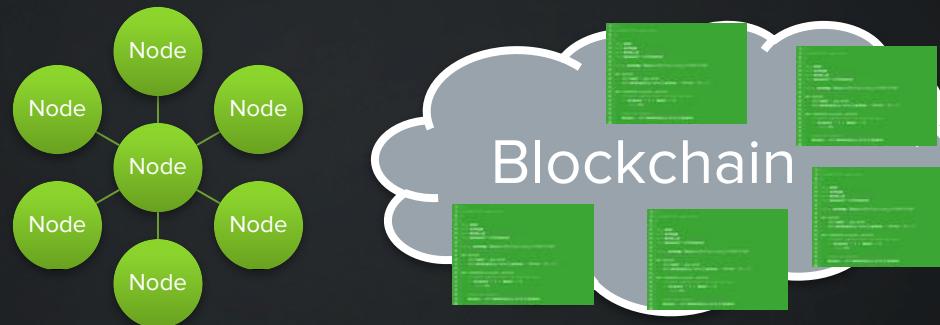
# Blockchain: Recap

- Network of nodes / peers
- They all maintain a copy of the Blockchain
- The Blockchain consists of blocks which store all network transactions
- Consensus between nodes is kept by a consensus algorithm (e.g. Proof of Work, Proof Of Stake, PBFT, ...)



# Smart Contracts

- The original idea (as published within the [Bitcoin paper](#)) included only **basic transactions within blocks**
- However, recently, Blockchain networks have appeared that support **storing programs** in the blockchain



- Such programs are called **smart contracts**, as they define **rules to alter the state** of the blockchain (e.g. variables stored in **on-chain-memory**)
- Each **node** within the network **verifies correctness** of program execution, hence the Blockchain is transformed from one **consistent** state to the next



# Chains with Smart Contract Support

Hyperledger:

- Suited for enterprise Blockchain solutions
- Open-source project started by Linux-Foundation



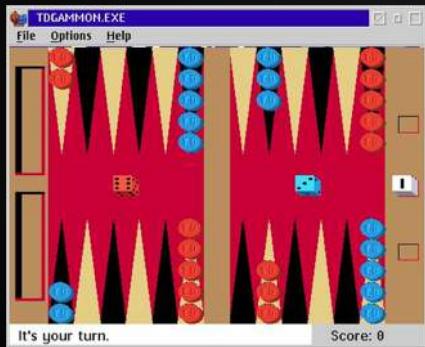
Ethereum:

- Blockchain implementation with smart contract support
- Easy tools for development of smart contracts and testing
- Good way to start and learn smart contract development

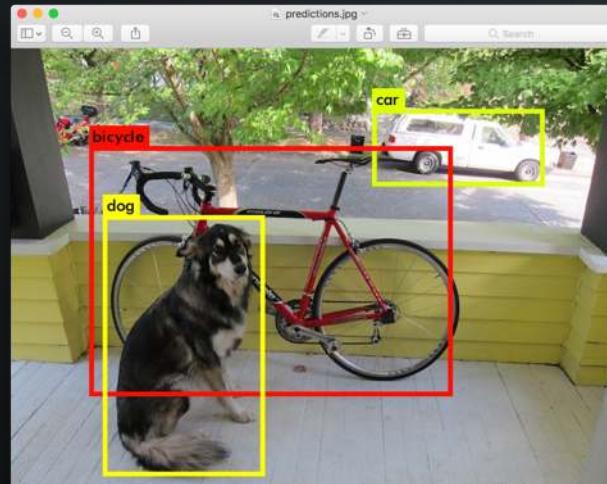


# Intro: Machine Learning

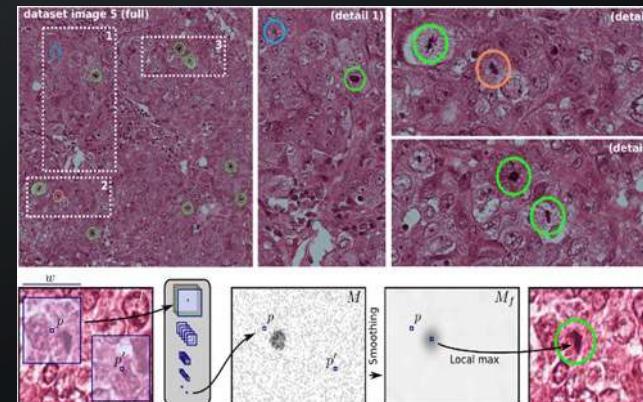
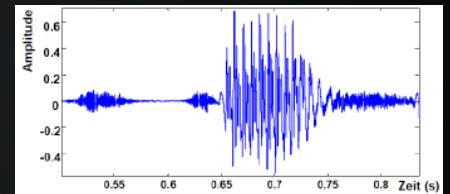
- **Machine learning** is the subfield of computer science that gives computers the ability to learn without being explicitly programmed. (Arthur Samuel, 1959).
- **Fantastic possibilities:** image recognition, text mining, time series prediction, anomaly detection, image de-noising, language translation, speech recognition, Backgammon, autonomous driving, and a lot more



Source:  
<https://en.wikipedia.org/wiki/TD-Gammon>



Source: <https://pjreddie.com/darknet/yolo/>



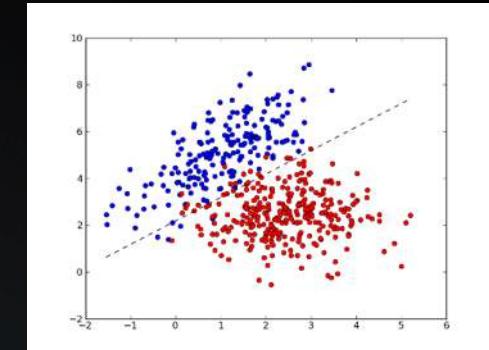
Source:  
<http://people.idsia.ch/~juergen/deeplearningwinsMICCAIgrandchallenge.html>



# Overview: Machine Learning

## Supervised Learning:

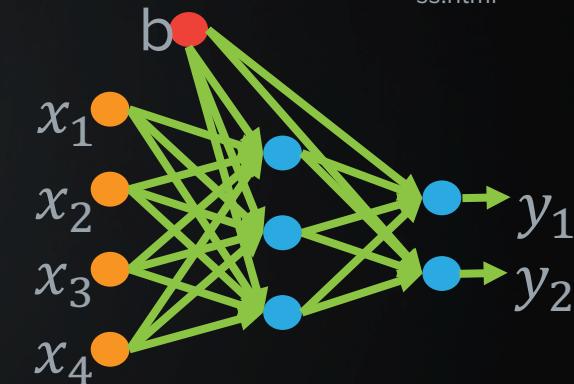
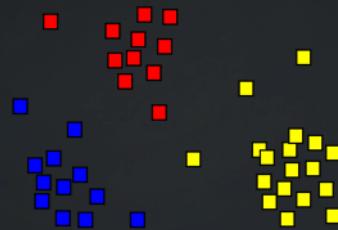
- Classification / Regression
- For each **input vector** a **label or target value** is given
- Algorithms: Multilayer Perceptron, Support Vector Machine, Random Forest, Adaboost, and many more



Source:  
[http://mipy.sourceforge.net/docs/3.5/lin\\_class.html](http://mipy.sourceforge.net/docs/3.5/lin_class.html)

## Unsupervised Learning:

- No labels / target given
- Purpose is to **find clusters** within data
- Algorithms: K-Means, DBSCAN, OPTICS, Expectation Maximization, Kohonen-Maps, and many more



## Reinforcement Learning:

- No exact label given, algorithm learns using **positive and negative reward**

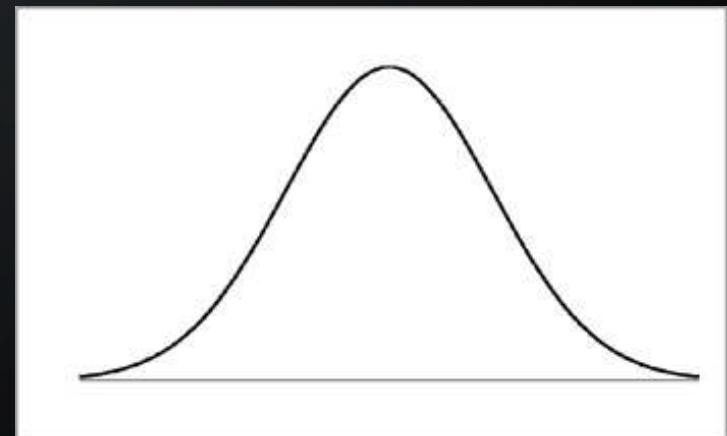
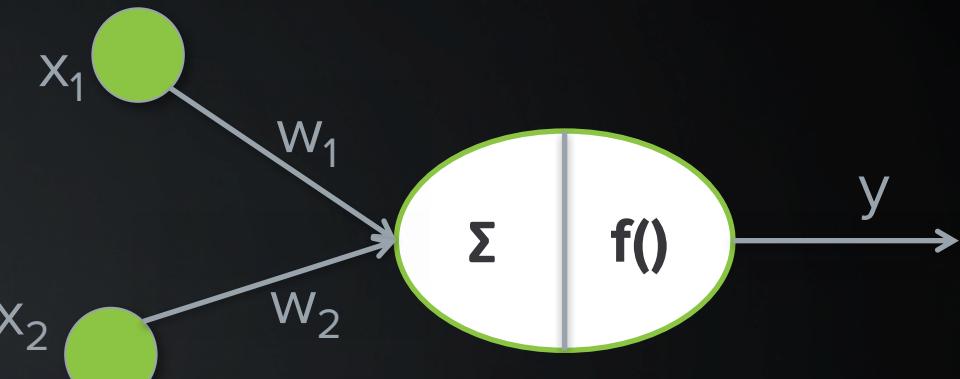
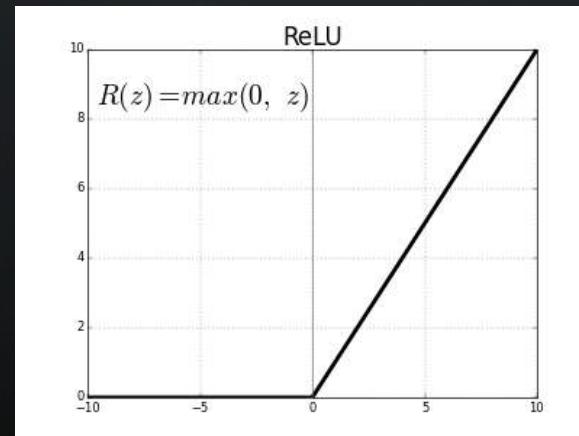
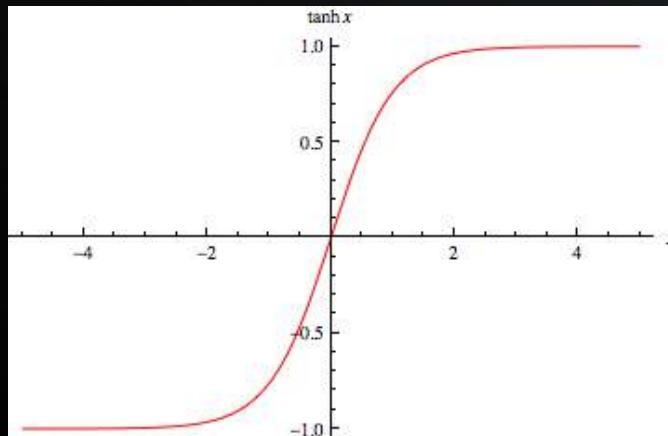


# Artificial Neural Networks (ANN)

- Biologically inspired family of algorithms to solve learning tasks
- We'll focus on supervised learning here

Neuron:

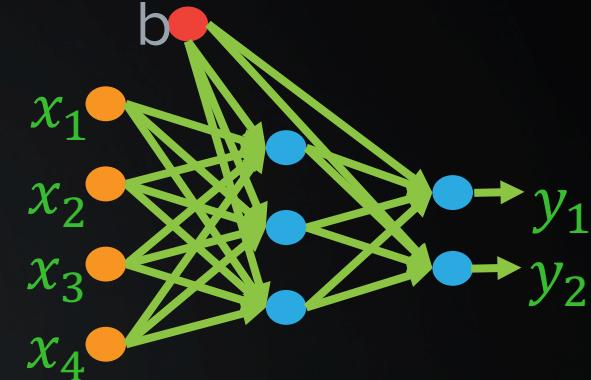
- Basic building-block of an ANN
- Maps several inputs to one output
- Each neuron has weights associated to its inputs
- $y = f(x_1w_1 + x_2w_2 + \dots + x_nw_n + b)$
- $f(z)$  could be: sigmoid:  $\tanh(z)$ , relu:  $\max(0, z)$ , linear, gaussian, etc.



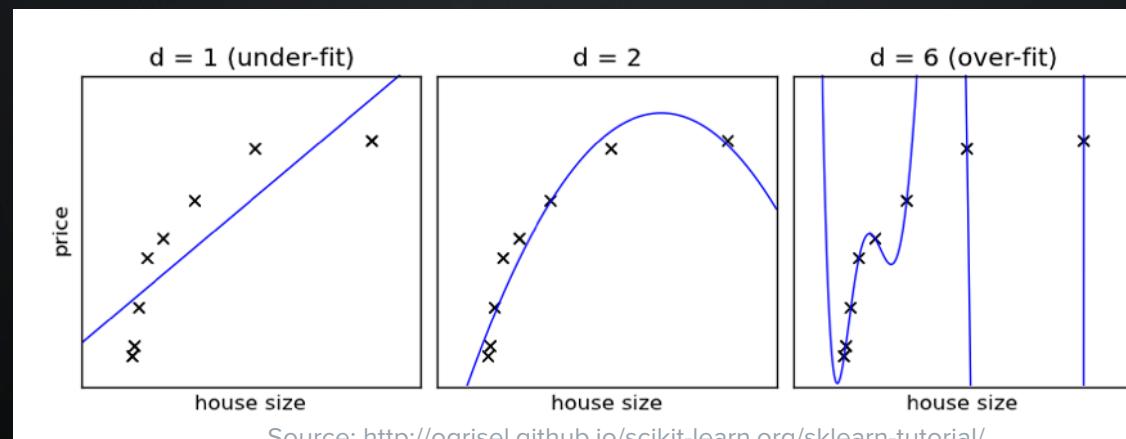
# Neurons: Stacking them Together

- Neurons can be combined into **multiple layers** forming a **simple ANN**:

Input vector  $\underline{x}$   
Output vector  $\mathbf{y}$



- The architecture will depend on the specific use-case, however:
- The **more neurons / weights** and layers within the network, the **more complex** the problems the network can solve
- **However:** more weights need more training data and careful handling:  
**overfitting**

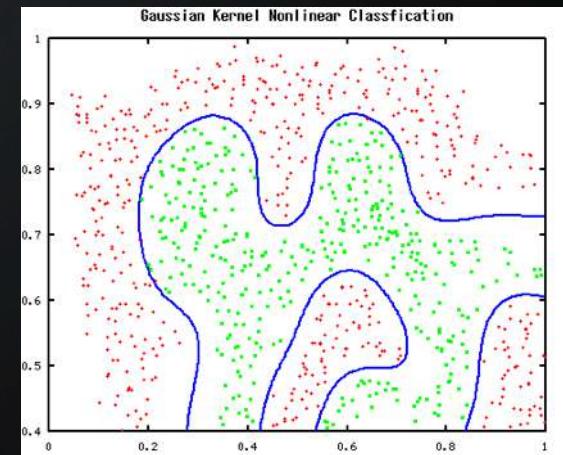


Source: <http://ogrissel.github.io/scikit-learn.org/sklearn-tutorial/>



# Training ANNs

- Objective: The network should be able to **generalize** and give a good prediction for **unknown input vectors**
- What can be **adjusted** within the network?
- → **Neuron weights!**
- We are searching for a **set of weights** that enables our network to **make accurate predictions**, that is, to give a **correct output** (category/value) for given inputs
- As this can be formulated as an **optimization problem**, ANNs are typically trained using optimization techniques, such as **error-backpropagation** (which is actually a gradient descent method)
- Objective: minimizing prediction / classification error

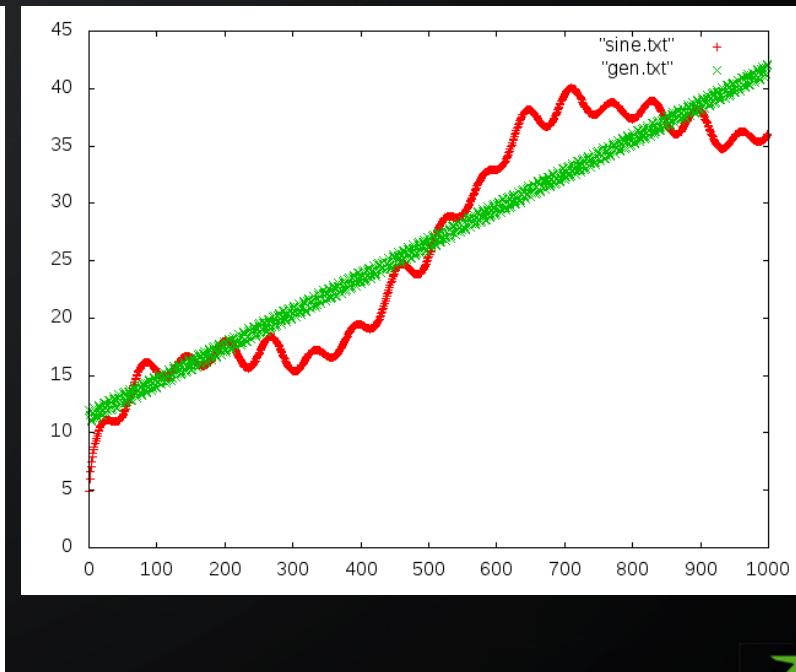
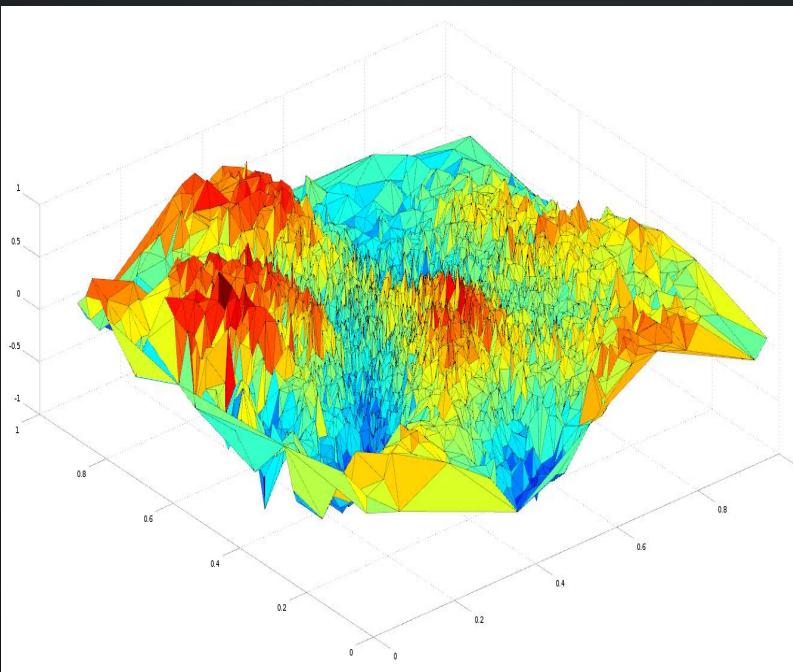
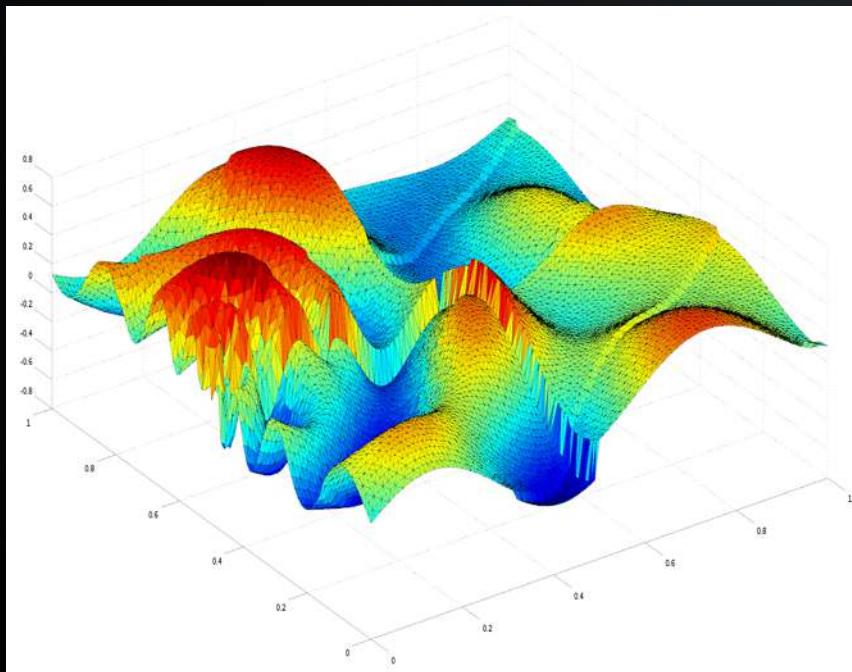


Source:  
<http://xenon.stanford.edu/~jianzh/ml/kernel.PNG>



# Some Training Details

- A batch of training samples (input vector  $x$  and correct output „teacher“ value  $t$  or class label) is fed into the network one by one
- Deviation between output vector  $y$  and  $t \rightarrow$  prediction error  $E$
- Training will try to find a minimum in the error surface over all training samples by adjusting neuron weights



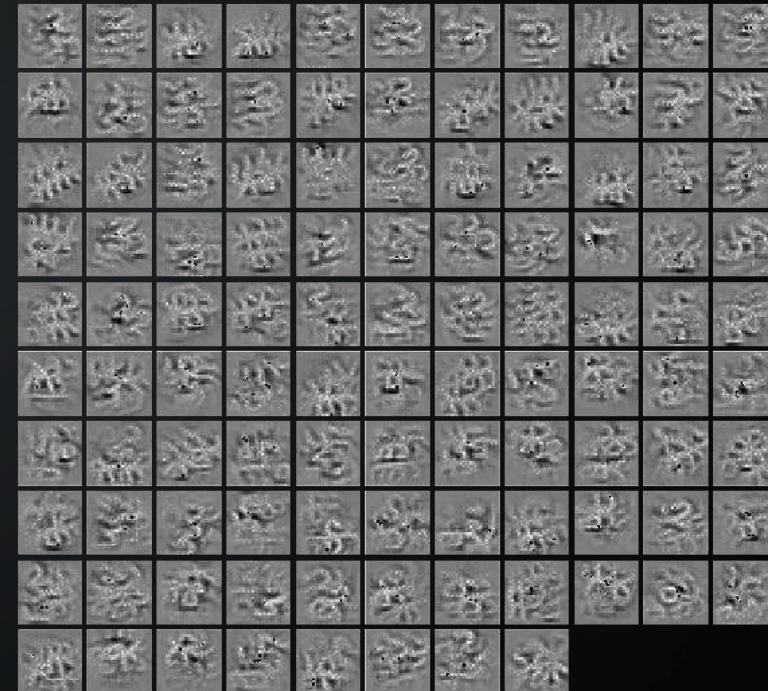
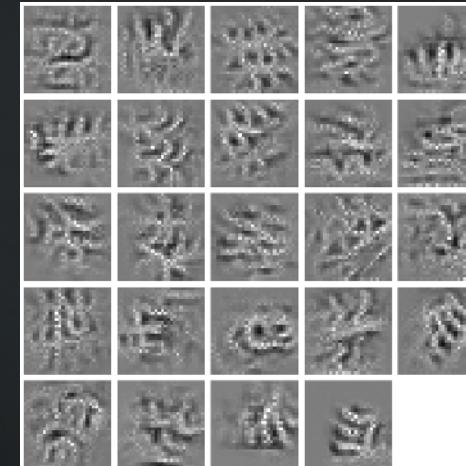
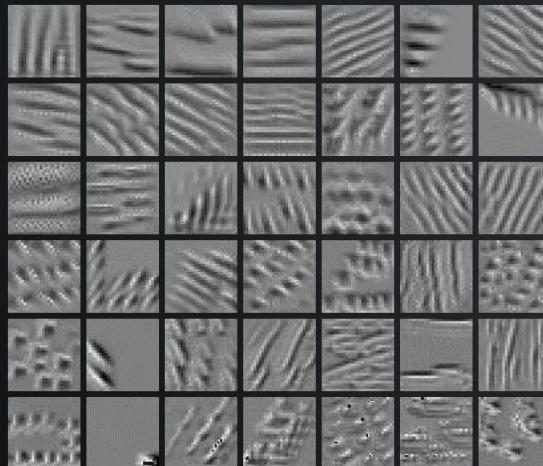
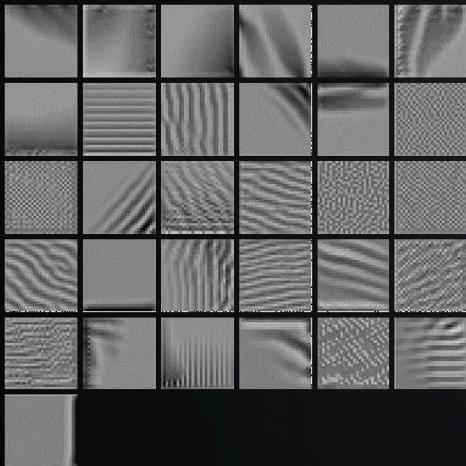
I need your help:

Let's beat the image classifier!



# Deep Learning: Motivation

- In recent years, computing power has increased dramatically
- Ever-growing amount of digital data that can be used for training
- Researchers found out that adding more hidden layers to an ANN will enable it to learn hierarchical abstraction levels of the data



- Important difficulties: vanishing / exploding gradients, overfitting
- Solutions: Relu activation, dropout, batch normalization

Source: Aufarbeitung aktueller Deep-Learning-Techniken  
– Daniel Seichter 2016 TU Ilmenau, NIKR

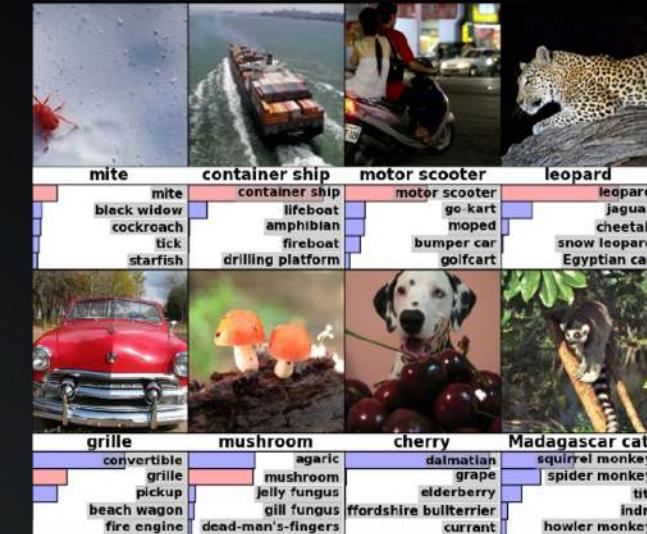


# Breaking Records

- Deep Learning is breaking records in many disciplines nowadays

Examples:

- 2009: Speech Recognition: A system based on Deep Learning could transform speech to text better than previous algorithms → Today Apple Siri and Android speech recognition both utilize DL
- 2012: Large Scale Visual Recognition Challenge 2012: Team around Geoffrey Hinton (Univ. Toronto, Google) introduced AlexNet and it had a top5-misclassification rate of only 15.3% (the winner one year before: 26,2%)
- 2016: Google DeepMind's AlphaGo won against Lee Sedol, one of the best Go-players world-wide (it uses only 1920 CPUs and 280 GPUs, you see Go is difficult ;-))



Source: A. Krizhevsky, et al.: Imagenet classification with deep convolutional neural networks, NIPS, 2012



Source: <http://www.latimes.com/business/la-fi-alpha-go-victory-20160315-htmlstory.html>



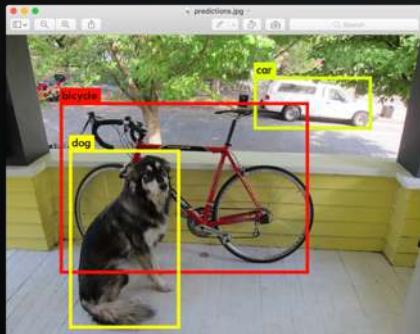
# Convolutional Neural Networks

Recap – Image Convolution:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 4 \\ 3 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 6 & 6 \\ 6 & 6 \end{bmatrix}$$

Step1:  $(1 * 1 + 2 * 0 + 4 * 0 + 5 * 1) = 6$

- Convolutional Neural Networks consist of **several convolutional layers**
- The next layer receives only the **filter responses** from the layer before
- Filter kernel **weights** are optimized during training as well

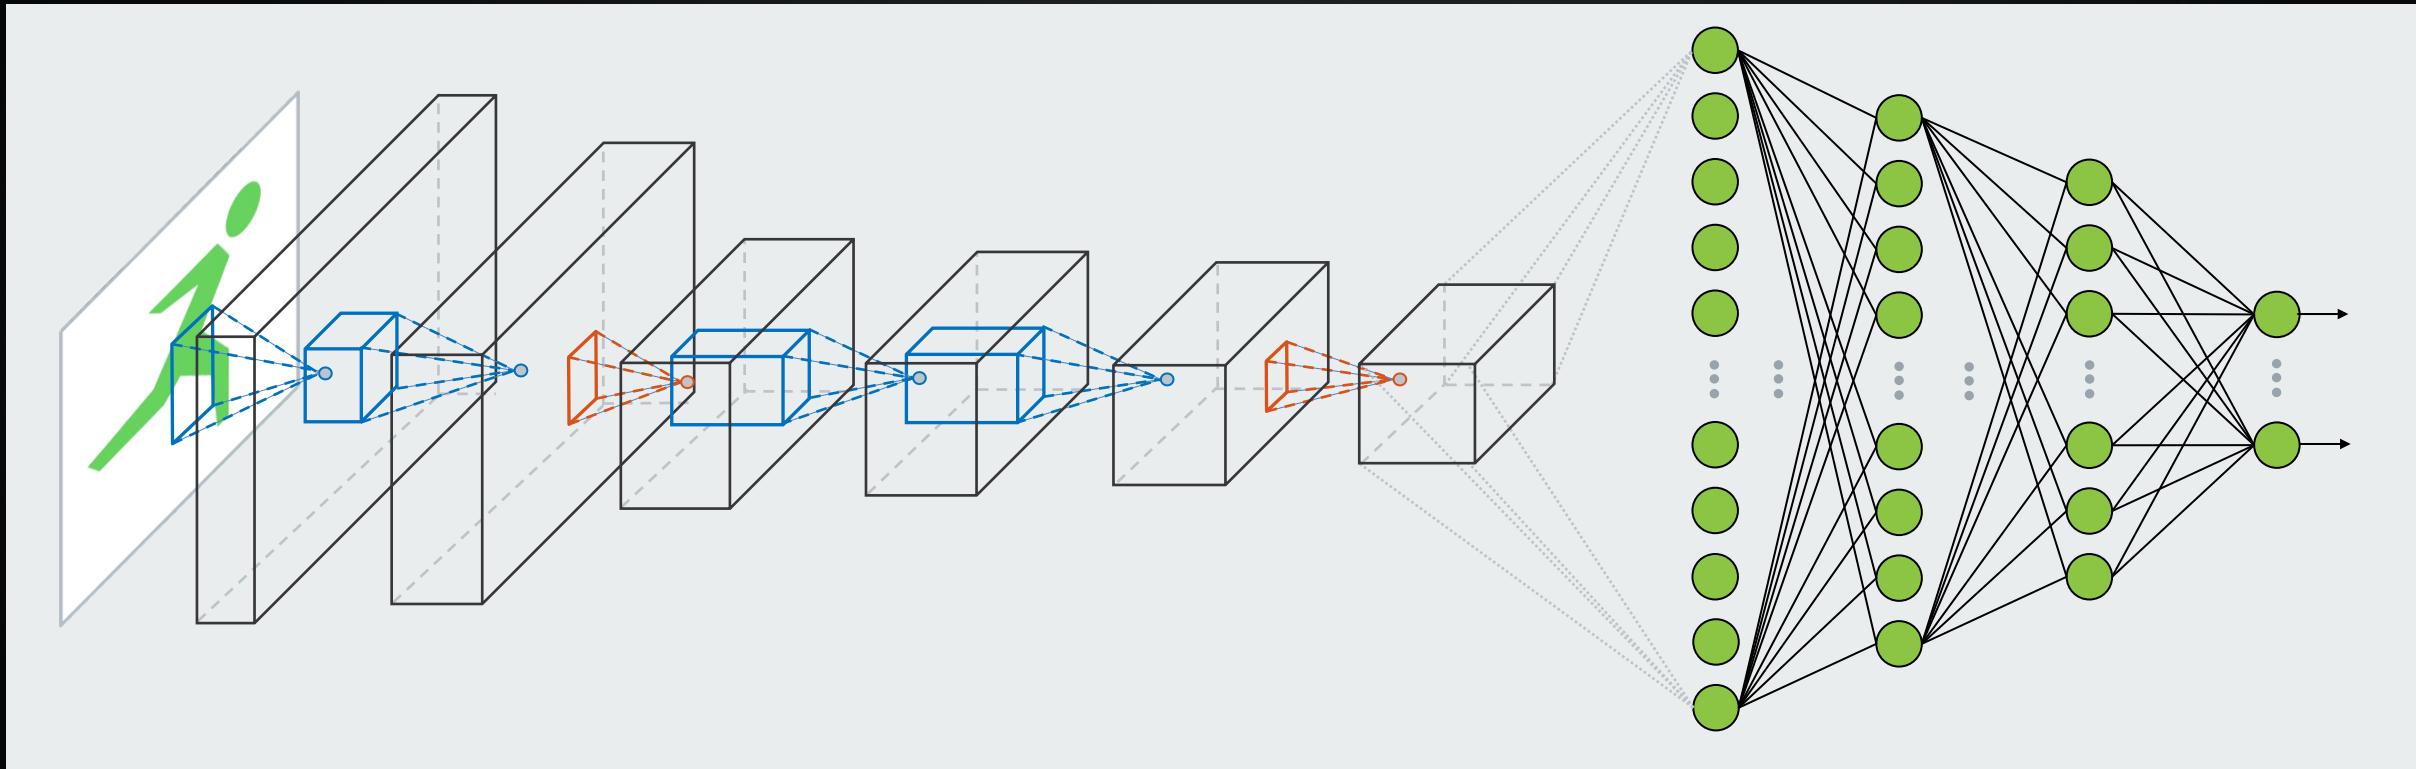


Source:  
<https://pjreddie.com/darknet/yolo/>



# Convolutional Networks [2]

Example Network Structure:

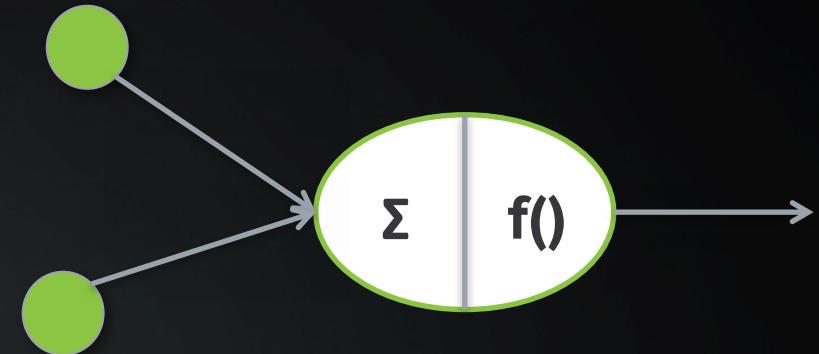


Source: Aufarbeitung aktueller Deep-Learning-Techniken – Daniel Seichter 2016 TU Ilmenau, NIKR



# Machine Learning: Recap

- Algorithms that **learn** to solve problems
- Artificial Neural Networks consist of **neurons**
- They can be used to **classify** input vectors (and for **regression** and **clustering**, etc.)
- Deep Neural Networks consist of **many hidden layers** and can **learn hierachical abstractions** within the input space
- Convolutional Neural Networks possess **convolutional layers** that apply trainable **filter kernels** on the data (e.g. images)



# Checkpoint

- Blockchain theory



- Machine Learning Overview



# Defining Our Objective

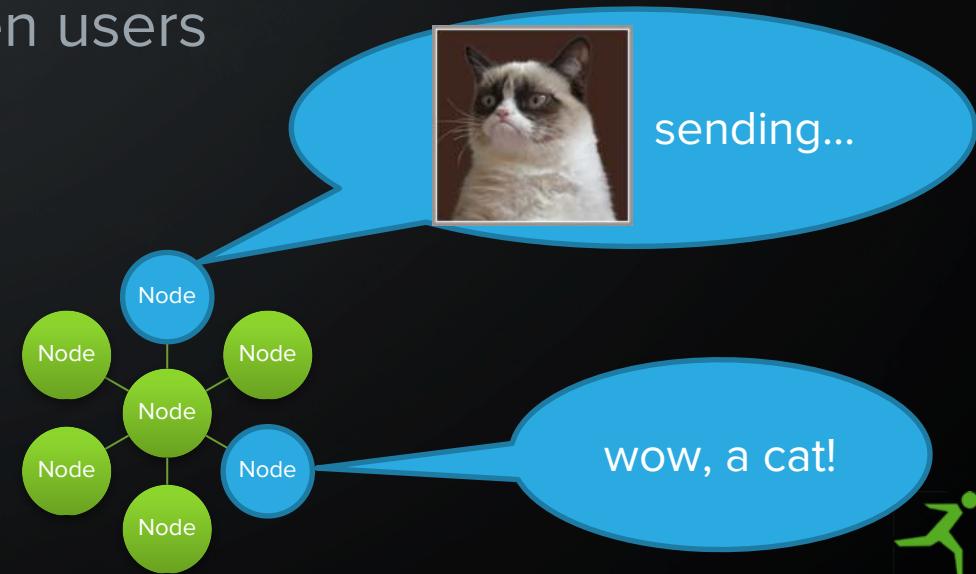
- Build a decentralized message service
- Send text messages and images to a blockchain
- Use Convolutional Neural Networks to extract image tags
- Subscribe to tags and receive the posts you're interested in
- Do all that in about 350 lines of Python code

```
--  
found contract on blockchain!  
--  
starting chat command line...  
>> send  
--  
[composing new message]  
message....: Hi, my name is Thomas  
image file.: img.jpg  
custom tags: #start #reply #test  
--  
sending...
```



# Why is that interesting? [1]

- Give you the building-blocks to create decentralized applications on your own
- The network is fully functional and tamper-proof as long as enough nodes are using it
- Data is timestamped, automatically analyzed, structured and easy to find
- Fast, intelligent and secure collaboration between users



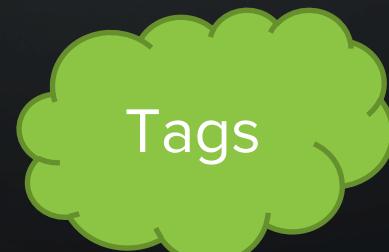
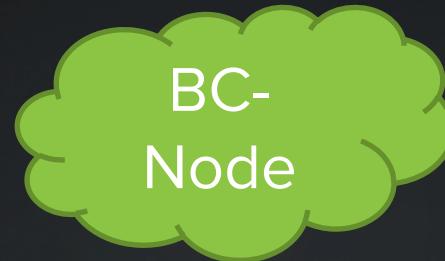
# Why is that interesting? [2]

- We can apply this technology to companies, e.g. for improved document and contract handling
- Documents can be analyzed for topics, images, reasonable recipients via Machine Learning
- Using a company-wide Blockchain, they will be recorded and timestamped, are immune to tampering and will be available in seconds
- By the way: we can do that for you ;-)

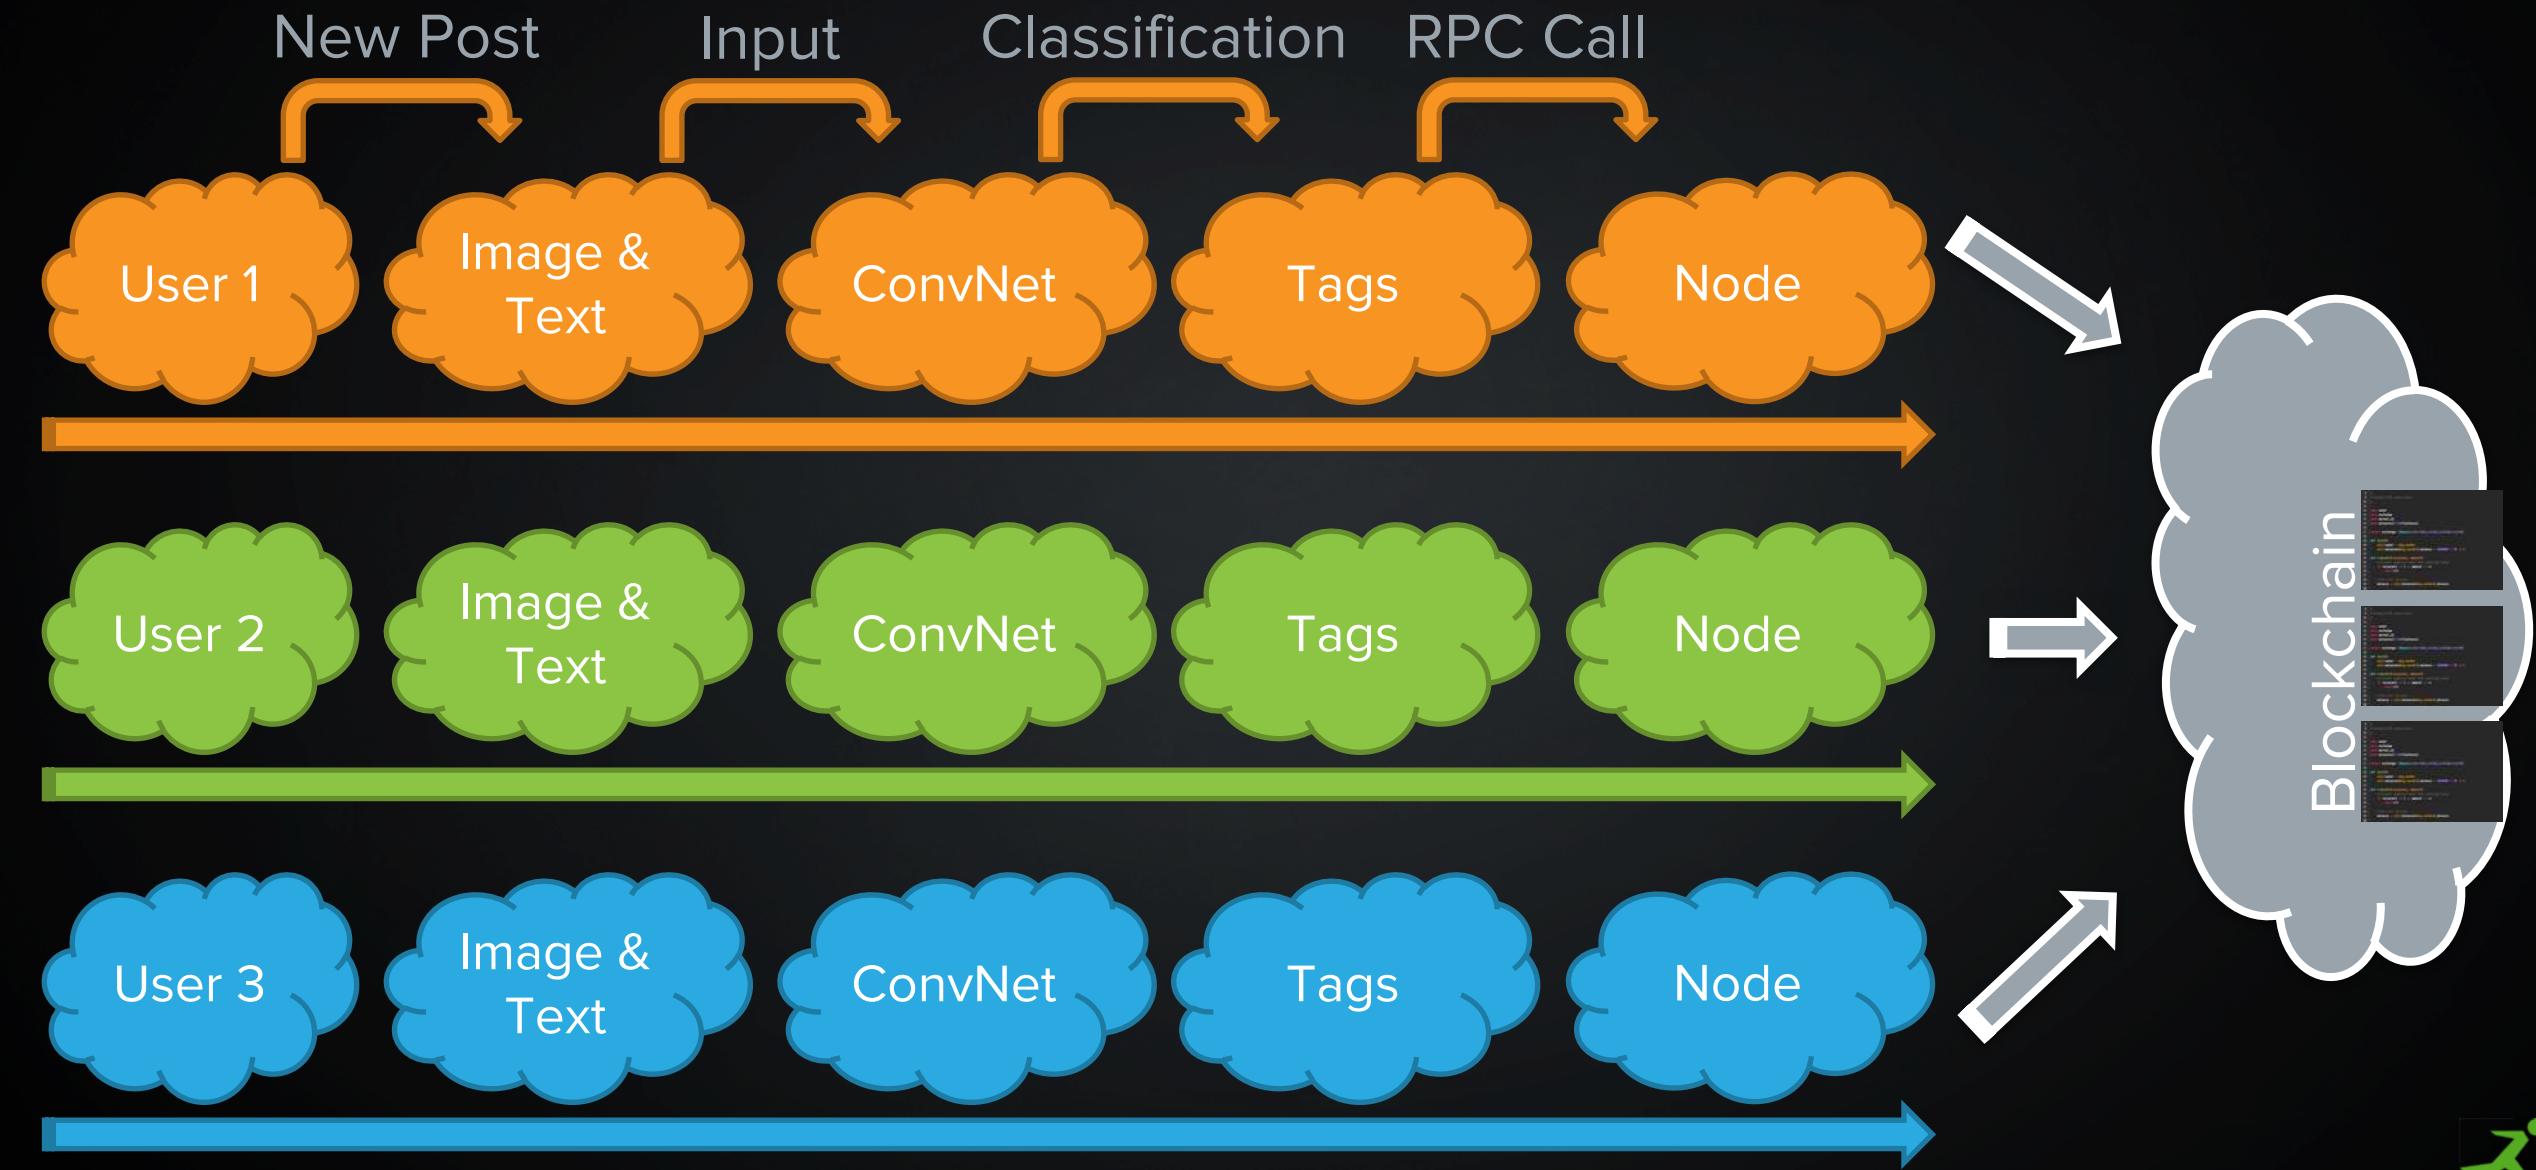


# Our Demo Architecture

How would you do it? Let's consider the following pieces:  
(Please fill in the diagram on your table, 5min)

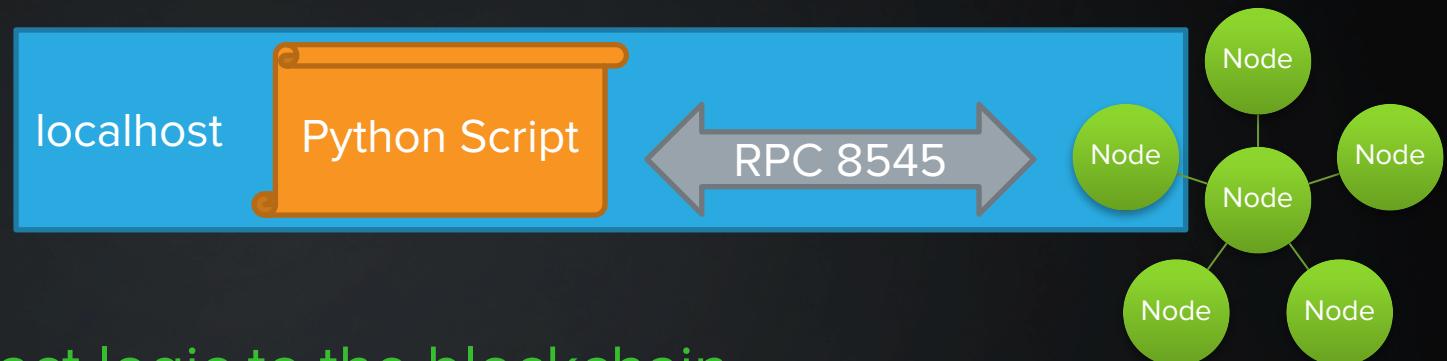


# Our Demo Architecture



# Our Demo Architecture

- Each user will run a **blockchain node**, together they will form the network
- Communication with the local blockchain node is done via **RPC** (remote procedure call)



- We need to **put our smart contract logic to the blockchain**
  - it contains a **mapping for every user account** to [message, image, tags]
  - for every new post, this content is changed, but the **history remains inside the chain**
- Finally, we need a **client tool for composing messages, subscribing to topics and receiving posts**



# Development Environment Setup

- We'll be using a **private development Ethereum blockchain** here (→ <https://github.com/ethereumjs>), since **testing and setup is very easy**

```
~/workspace/wrk/workshop$ testrpc
```

- Local testnet that **simulates real blockchain network**
- **Identical interface and behavior**
- **10 user accounts** automatically created
- Listening for **RPC requests** on port 8545

```
EthereumJS TestRPC v3.0.3

Available Accounts
=====
(0) 0x727ac9f9a3a470cd1437dc3cfb0bf55566b13e01
(1) 0x51a0039aa4ef8b2408bbdbcda83f53deb42be874
(2) 0x09e956ffc8fba560befc175eeeca55ee26542cffa
(3) 0x24d3d62dad12a8dd27ad3532ccb5a1c005a3eafe
(4) 0x5ca962eecae2a50bb7679277fcfd10faddac81a3
(5) 0x059f0391ea6102e565559110a05001db9503f099
(6) 0x2e9203da725ba08877925e3cf8dad28df9ed2844
(7) 0x9f2ef7919af6442b66afcd281cd7422653d0bf923
(8) 0xa0cb7e68bd6744c4fd6916846ce3bb4e3e3fb4ca
(9) 0x0c1202e576f83fe09ab1584212eddcea80d1b6f7

Private Keys
=====
(0) 8296df3e1b8f5f42a67503c779efd385b234c416d8578ce40bf9a4e2cec4b9a7
(1) 757d480a9425422de317f03dceec554ecc516f671df265bc697624631ecde9da
(2) f81fd407a79744a457ae7d0bc433ab0d1ac69909d25f729b6fff4be3f559a443
(3) 1658310940d424ba5437c343b4fe09a998c8d48ed4278fc30d5f0db180de9250
(4) 48a681c6d34c7cd3b98454f8c5fd35bb16d24a2e6a0d0705d87c598d339ae429
(5) 61cafbee004113a6e1dc27b0e9b96c251cf48bbb9ccdd0b38c7bf11f35d3f29d
(6) e1fc0c0df432d6081ee88ba1b2210288e8ed878f2296be83c965571c14df2b27
(7) 60a26e44fb7fcf7911993bb920229679e539b40647d8c5ec7651ce220be999da
(8) 8ebab6fa63fad7c65ab74a0a25cff49a0bfef7b20abbcc774d90a2a6749f75d3
(9) 5245e9e7b9183fc59f87c6ed7e9abd886ce8201ec72ed9858701319c89fc17bf

HD Wallet
=====
Mnemonic:      girl cry suggest minor valve orient mouse elder evolve animal puppy pumpkin
Base HD Path: m/44'/60'/0'/0/{account_index}

Listening on localhost:8545
```

# Package: Python

- We'll be using the **Python** programming language here, because it's cool (→ <https://www.python.org/>)



```
Python 2.7.12 (default, Nov 19 2016, 06:48:10)
[GCC 5.4.0 20160609] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> 
```



# Package: ethjsonrpc

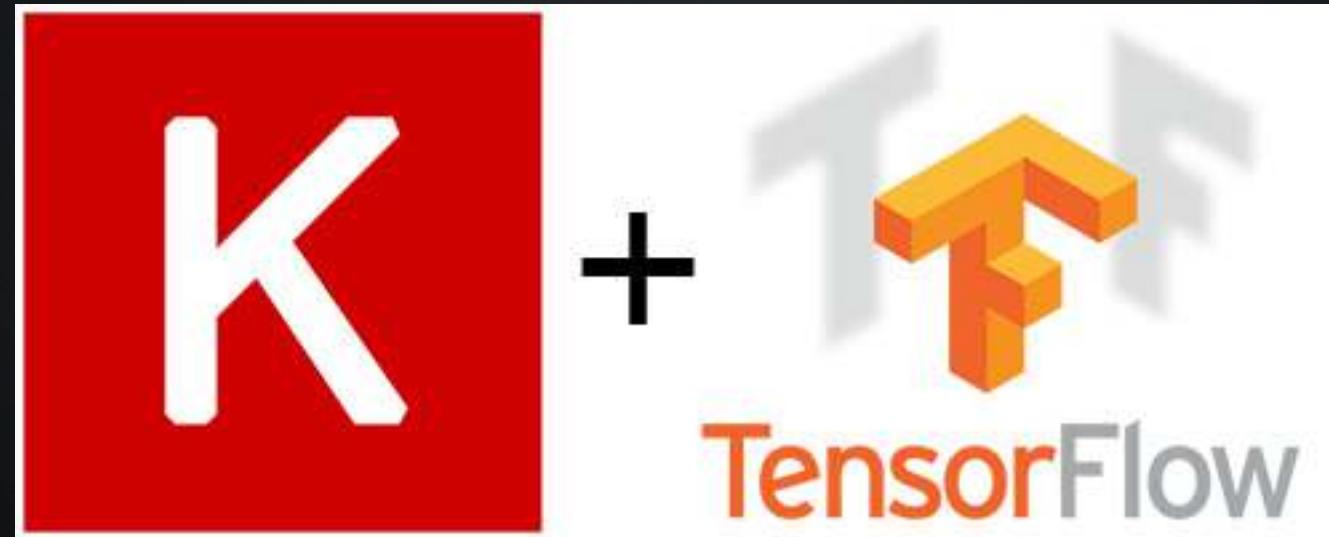
- For communication with the blockchain node, let's install the `ethjsonrpc` package
- It will provide an interface for communicating with our local blockchain node

```
Python 2.7.12 (default, Nov 19 2016, 06:48:10)
[GCC 5.4.0 20160609] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> from ethjsonrpc import EthJsonRpc
>>> rpc = EthJsonRpc('localhost', 8545)
>>> rpc.eth_blockNumber()
19
>>> rpc.eth_blockNumber()
19
>>> rpc.eth_blockNumber()
20
>>> rpc.eth_blockNumber()
20
>>> rpc.eth_blockNumber()
21
>>> █
```



# Packages: keras + tensorflow

- Keras with Tensorflow backend will take care of **image classification** for us, in **6 lines of code**



Now, let's get our hands dirty.  
Step-by-Step review of the client scripts:

Smart Contract

Deployment Script

Messaging Client



# That was just an Overview

- More Machine Learning stuff: Natural Language Processing, Document Topic Extraction, Optical Character Recognition, ...
- More Blockchain stuff: Smarter Smart Contracts, Blockchain Analytics (the chain has all the data = topics = messages = images, ...), ...
- Any questions? Ask them now, or drop me an e-mail. 

[t.schmiedel@reply.de](mailto:t.schmiedel@reply.de)

